

COMP3141

Software System Design and Implementation

Effects and State

Liam O'Connor
CSE, UNSW (and Data61)
Term 2 2019

Effects

Effects

Effects are observable phenomena from the execution of a program.

Effects

Effects

Effects are observable phenomena from the execution of a program.

Example (Memory effects)

```
int *p = ...  
... // read and write  
*p = *p + 1;
```

Effects

Effects

Effects are observable phenomena from the execution of a program.

Example (Memory effects)

```
int *p = ...  
... // read and write  
*p = *p + 1;
```

Example (IO)

```
// console IO  
c = getchar();  
printf("%d", 32);
```

Effects

Effects

Effects are observable phenomena from the execution of a program.

Example (Memory effects)

```
int *p = ...  
... // read and write  
*p = *p + 1;
```

Example (IO)

```
// console IO  
c = getchar();  
printf("%d", 32);
```

Example (Control flow)

```
// exception effect  
throw new Exception();
```

Effects

Effects

Effects are observable phenomena from the execution of a program.

Example (Memory effects)

```
int *p = ...  
... // read and write  
*p = *p + 1;
```

Example (IO)

```
// console IO  
c = getchar();  
printf("%d", 32);
```

Example (Non-termination)

```
// infinite loop  
while (1) {};
```

Example (Control flow)

```
// exception effect  
throw new Exception();
```

Internal vs. External Effects

External Observability

An *external* effect is an effect that is *observable* outside the function.

Internal effects are not observable from outside.

Internal vs. External Effects

External Observability

An *external* effect is an effect that is *observable* outside the function.

Internal effects are not observable from outside.

Example (External effects)

Console, file and network I/O; termination and non-termination; non-local control flow; etc.

Internal vs. External Effects

External Observability

An *external* effect is an effect that is *observable* outside the function.

Internal effects are not observable from outside.

Example (External effects)

Console, file and network I/O; termination and non-termination; non-local control flow; etc.

Are memory effects *external* or *internal*?

Internal vs. External Effects

External Observability

An *external* effect is an effect that is *observable* outside the function.

Internal effects are not observable from outside.

Example (External effects)

Console, file and network I/O; termination and non-termination; non-local control flow; etc.

Are memory effects *external* or *internal*?

Answer: Depends on the scope of the memory being accessed.
Global variable accesses are *external*.

Purity

A function with no external effects is called a *pure* function.

Pure functions

A *pure function* is the mathematical notion of a function. That is, a function of type $a \rightarrow b$ is *fully* specified by a mapping from all elements of the domain type a to the codomain type b .

Purity

A function with no external effects is called a *pure* function.

Pure functions

A *pure function* is the mathematical notion of a function. That is, a function of type $a \rightarrow b$ is *fully* specified by a mapping from all elements of the domain type a to the codomain type b .

Consequences:

- Two invocations with the same arguments result in the same value.

Purity

A function with no external effects is called a *pure* function.

Pure functions

A *pure function* is the mathematical notion of a function. That is, a function of type $a \rightarrow b$ is *fully* specified by a mapping from all elements of the domain type a to the codomain type b .

Consequences:

- Two invocations with the same arguments result in the same value.
- No observable trace is left beyond the result of the function.

Purity

A function with no external effects is called a *pure* function.

Pure functions

A *pure function* is the mathematical notion of a function. That is, a function of type $a \rightarrow b$ is *fully* specified by a mapping from all elements of the domain type a to the codomain type b .

Consequences:

- Two invocations with the same arguments result in the same value.
- No observable trace is left beyond the result of the function.
- No implicit notion of time or order of execution.

Purity

A function with no external effects is called a *pure* function.

Pure functions

A *pure function* is the mathematical notion of a function. That is, a function of type $a \rightarrow b$ is *fully* specified by a mapping from all elements of the domain type a to the codomain type b .

Consequences:

- Two invocations with the same arguments result in the same value.
- No observable trace is left beyond the result of the function.
- No implicit notion of time or order of execution.

Question: Are Haskell functions *pure*?

Haskell Functions

Haskell functions are technically **not** pure.

Haskell Functions

Haskell functions are technically **not** pure.

- They can loop infinitely.

Haskell Functions

Haskell functions are technically **not** pure.

- They can loop infinitely.
- They can throw exceptions (**partial functions**).

Haskell Functions

Haskell functions are technically **not** pure.

- They can loop infinitely.
- They can throw exceptions (**partial functions**).
- They can force evaluation of unevaluated expressions.

Haskell Functions

Haskell functions are technically **not** pure.

- They can loop infinitely.
- They can throw exceptions (**partial functions**).
- They can force evaluation of unevaluated expressions.

Caveat

Purity only applies to a particular level of abstraction. Even ignoring the above, assembly instructions produced by GHC aren't really pure.

Haskell Functions

Haskell functions are technically **not** pure.

- They can loop infinitely.
- They can throw exceptions (**partial functions**).
- They can force evaluation of unevaluated expressions.

Caveat

Purity only applies to a particular level of abstraction. Even ignoring the above, assembly instructions produced by GHC aren't really pure.

Despite the impurity of Haskell functions, we can often reason as though they are pure. Hence we call Haskell a **purely functional** language.

The Danger of Implicit Side Effects

- They introduce (often subtle) requirements on the evaluation order.

The Danger of Implicit Side Effects

- They introduce (often subtle) requirements on the evaluation order.
- They are not visible from the type signature of the function.

The Danger of Implicit Side Effects

- They introduce (often subtle) requirements on the evaluation order.
- They are not visible from the type signature of the function.
- They introduce **non-local** dependencies which is bad for software design, increasing *coupling*.

The Danger of Implicit Side Effects

- They introduce (often subtle) requirements on the evaluation order.
- They are not visible from the type signature of the function.
- They introduce **non-local** dependencies which is bad for software design, increasing *coupling*.
- They interfere badly with strong typing, for example mutable arrays in Java, or reference types in ML.

The Danger of Implicit Side Effects

- They introduce (often subtle) requirements on the evaluation order.
- They are not visible from the type signature of the function.
- They introduce **non-local** dependencies which is bad for software design, increasing *coupling*.
- They interfere badly with strong typing, for example mutable arrays in Java, or reference types in ML.

We can't, in general, **reason equationally** about effectful programs!

Can we program with pure functions?

Yes! We've been doing it for the past 6 weeks.

Can we program with pure functions?

Yes! We've been doing it for the past 6 weeks.

Typically, a computation involving some state of type s and returning a result of type a can be expressed as a function:

$$s \rightarrow (s, a)$$

Can we program with pure functions?

Yes! We've been doing it for the past 6 weeks.

Typically, a computation involving some state of type s and returning a result of type a can be expressed as a function:

$$s \rightarrow (s, a)$$

Rather than **change** the state, we return a **new copy** of the state.

Can we program with pure functions?

Yes! We've been doing it for the past 6 weeks.

Typically, a computation involving some state of type s and returning a result of type a can be expressed as a function:

$$s \rightarrow (s, a)$$

Rather than **change** the state, we return a **new copy** of the state.

Efficiency?

All that copying might seem expensive, but by using tree data structures, we can usually reduce the cost to an $\mathcal{O}(\log n)$ overhead.

State Passing

Example (Labelling Nodes)

```
data Tree a = Branch a (Tree a) (Tree a) | Leaf
```


State Passing

Example (Labelling Nodes)

```
data Tree a = Branch a (Tree a) (Tree a) | Leaf
```

Given a tree, label each node with an ascending number in infix order:

```
label :: Tree () -> Tree Int
```




State Passing

Example (Labelling Nodes)

```
data Tree a = Branch a (Tree a) (Tree a) | Leaf
```

Given a tree, label each node with an ascending number in infix order:

```
label :: Tree () -> Tree Int
```




State Passing

Example (Labelling Nodes)

```
data Tree a = Branch a (Tree a) (Tree a) | Leaf
```

Given a tree, label each node with an ascending number in infix order:

```
label :: Tree () -> Tree Int
```



Let's use **monads** to simplify this!

State Monads

```
newtype State s a = State (s -> (s, a))
```

State Monads

```
newtype State s a = State (s -> (s, a))
```

State Monad

```
get :: State s s  
put :: s -> State s ()  
modify :: (s -> s) -> State s ()
```

State Monads

```
newtype State s a = State (s -> (s, a))
```

State Monad

```
get  :: State s s  
put  :: s -> State s ()  
modify :: (s -> s) -> State s ()
```

Here we use a **monadic** interface to simplify the passing of our state around, so that we don't need to manually plumb data around.

Effects

Sometimes we need side effects.

- We need to perform I/O, to communicate with the user or hardware.
- We might need effects for maximum efficiency.
(but usually internal effects are sufficient)

Effects

Sometimes we need side effects.

- We need to perform I/O, to communicate with the user or hardware.
- We might need effects for maximum efficiency.
(but usually internal effects are sufficient)

Haskell's approach

Pure by default. Effectful when necessary.

The IO Type

A **procedure** that performs some side effects, returning a result of type `a` is written as `IO a`.

The IO Type

A **procedure** that performs some side effects, returning a result of type `a` is written as `IO a`.

World interpretation

`IO a` is an abstract type. But we can think of it as a function:

$$\text{RealWorld} \rightarrow (\text{RealWorld}, a)$$

(that's how it's implemented in GHC)

The IO Type

A **procedure** that performs some side effects, returning a result of type `a` is written as `IO a`.

World interpretation

`IO a` is an abstract type. But we can think of it as a function:

$$\text{RealWorld} \rightarrow (\text{RealWorld}, a)$$

(that's how it's implemented in GHC)

```
(>>=) :: IO a -> (a -> IO b) -> IO b
```

```
pure  :: a -> IO a
```

```
getChar :: IO Char
```

```
readLine :: IO String
```

```
putStrLn :: String -> IO ()
```

Infectious IO

We can convert pure values to impure procedures with pure:

`pure :: a -> IO a`

Infectious IO

We can convert pure values to impure procedures with `pure`:

```
pure :: a -> IO a
```

But we can't convert impure procedures to pure values:

```
???? :: IO a -> a
```

Infectious IO

We can convert pure values to impure procedures with `pure`:

```
pure :: a -> IO a
```

But we can't convert impure procedures to pure values:

```
???? :: IO a -> a
```

The only function that gets an `a` from an `IO a` is `>>=`:

```
(>>=) :: IO a -> (a -> IO b) -> IO b
```

But it returns an `IO` procedure as well.

Infectious IO

We can convert pure values to impure procedures with pure:

```
pure :: a -> IO a
```

But we can't convert impure procedures to pure values:

```
???? :: IO a -> a
```

The only function that gets an `a` from an `IO a` is `>>=`:

```
(>>=) :: IO a -> (a -> IO b) -> IO b
```

But it returns an `IO` procedure as well.

Conclusion

The moment you use an `IO` procedure in a function, `IO` shows up in the types, and you can't get rid of it!

If a function makes use of `IO` effects directly or indirectly, it will have `IO` in its type!

Haskell Design Strategy


We ultimately “run” IO procedures by calling them from `main`:

```
main :: IO ()
```

Haskell Design Strategy

We ultimately “run” IO procedures by calling them from `main`:

```
main :: IO ()
```



Examples

Example (Triangles)

Given an input number n , print a triangle of $*$ characters of base width n .

Examples

Example (Triangles)

Given an input number n , print a triangle of $*$ characters of base width n .

Example (Maze Game)

Design a game that reads in a $n \times n$ maze from a file. The player starts at position $(0, 0)$ and must reach position $(n - 1, n - 1)$ to win. The game accepts keyboard input to move the player around the maze.

Benefits of an IO Type

- Absence of effects makes type system more informative:
 - A type signatures captures **entire interface** of the function.
 - All **dependencies are explicit** in the form of data dependencies.
 - All **dependencies are typed**.

Benefits of an IO Type

- Absence of effects makes type system more informative:
 - A type signatures captures **entire interface** of the function.
 - All **dependencies are explicit** in the form of data dependencies.
 - All **dependencies are typed**.
- It is easier to reason about pure code and it is easier to test:
 - Testing is local, doesn't require complex set-up and tear-down.
 - Reasoning is local, doesn't require state invariants.
 - Type checking leads to strong guarantees.

Mutable Variables

We can have honest-to-goodness mutability in Haskell, if we really need it, using `IORef`.

```
data IORef a
newIORef :: a -> IO (IORef a)
readIORef :: IORef a -> IO a
writeIORef :: IORef a -> a -> IO ()
```

Mutable Variables

We can have honest-to-goodness mutability in Haskell, if we really need it, using `IORef`.

```
data IORef a
newIORef :: a -> IO (IORef a)
readIORef :: IORef a -> IO a
writeIORef :: IORef a -> a -> IO ()
```

Example (Effectful Average)

Average a list of numbers using `IORefs`.

Mutable Variables, Locally

Something like averaging a list of numbers doesn't require external effects, even if we use mutation internally.

Mutable Variables, Locally

Something like averaging a list of numbers doesn't require external effects, even if we use mutation internally.

```
data STRef s a
newSTRef :: a -> ST (STRef s a)
readSTRef :: STRef s a -> ST s a
writeSTRef :: STRef s a -> a -> ST s ()
runST :: (forall s. ST s a) -> a
```


Mutable Variables, Locally

Something like averaging a list of numbers doesn't require external effects, even if we use mutation internally.

```
data STRef s a
newSTRef :: a -> ST (STRef s a)
readSTRef :: STRef s a -> ST s a
writeSTRef :: STRef s a -> a -> ST s ()
runST :: (forall s. ST s a) -> a
```

The extra *s* parameter is called a **state thread**, that ensures that mutable variables don't leak outside of the ST computation.

Mutable Variables, Locally

Something like averaging a list of numbers doesn't require external effects, even if we use mutation internally.

```
data STRef s a
newSTRef :: a -> ST (STRef s a)
readSTRef :: STRef s a -> ST s a
writeSTRef :: STRef s a -> a -> ST s ()
runST :: (forall s. ST s a) -> a
```

The extra *s* parameter is called a **state thread**, that ensures that mutable variables don't leak outside of the ST computation.

Note

The ST Monad is not assessable in this course, but it is useful sometimes in Haskell programming.

QuickChecking Monads

QuickCheck lets us test IO (and ST) using this special **property monad** interface:

```
monadicIO :: PropertyM IO () -> Property
pre       :: Bool -> PropertyM IO ()
assert    :: Bool -> PropertyM IO ()
run       :: IO a -> PropertyM IO a
```

QuickChecking Monads

QuickCheck lets us test IO (and ST) using this special **property monad** interface:

```
monadicIO :: PropertyM IO () -> Property
pre       :: Bool -> PropertyM IO ()
assert    :: Bool -> PropertyM IO ()
run       :: IO a -> PropertyM IO a
```

Example (Testing average)

Let's test that our IO average function works like the non-effectful one.

QuickChecking Monads

QuickCheck lets us test IO (and ST) using this special **property monad** interface:

```
monadicIO :: PropertyM IO () -> Property
pre       :: Bool -> PropertyM IO ()
assert    :: Bool -> PropertyM IO ()
run       :: IO a -> PropertyM IO a
```

Example (Testing average)

Let's test that our IO average function works like the non-effectful one.

Example (Testing gfactor)

Let's test that the GNU factor program works correctly!

Homework

- 1 New exercise out, due Tuesday next week.
- 2 Last week's quiz is due on Friday.
- 3 This week's quiz is due the following Friday.
- 4 **Note:** Assignment 2 released next week!